

MSc in Computer Science
at University of Milan

Statistical Methods for Machine Learning
Kernelized Linear Predictor
course held by **Nicoló Cesa-Bianchi**

Email:
davide.cognigni@studenti.unimi.it

Created by:
Davide Cognigni
mat. 09732A

Academic year of 2023/2024

Contents

1	introduction	4
1.1	Project Description	4
1.2	Project Structure	4
1.3	Usage	4
2	Dataset analysis and preprocessing	6
2.1	Dataset description	6
2.2	Feature Scaling	6
2.3	Outliers removal	7
2.4	Feature correlation	7
2.5	Feature expansion	8
3	Perceptron	9
3.1	Naive version	10
3.2	Feature Expansion of 2nd degree	10
4	Support Vector Machine	12
4.1	Naive	13
4.1.1	implementation details	13
4.1.2	Hyperparameter tuning	14
4.2	Logistic regression	15
4.3	Feature Expansion	15
4.3.1	Pegasos	15
4.3.2	Logistic Regression	16
5	Kernel	17

1 introduction

1.1 Project Description

The goal of this project is to train different learning algorithms to solve the binary classification problem. Specifically, we need to predict the value of the label in the y column based on the numerical features x_1 through x_{10} .

As mentioned in the description of the project, I used the zero-one loss as a metric to evaluate the performance of the algorithms.

In the following chapter I describe how I analyzed and preprocessed the dataset to improve the performance of the modules.

In the following chapters I describe how I implemented each algorithm and how I chose the hyperparameters for them.

I have also analyzed how varying the hyperparameter affects the training and testing error and explained these empirical results with the theoretical background provided by the lectures.

I implemented the following algorithms:

- Perceptron
- Pegasos
- Logistic Pegasos
- Feature expanded Perceptron (with 2nd degree polynomial expansion)
- Feature expanded Pegasos (with 2nd degree polynomial expansion)
- Feature expanded Pegasos with logistic loss (with 2nd degree polynomial expansion)
- Kernel Perceptron
- Kernel Pegasos

1.2 Project Structure

The project is divided into the following folders:

- datasets: contains the provided dataset
- models: a set of pre-trained models, created using the train subcommand
- src: the source code of the project, the entry point is main.py and can be used both to train the models from the dataset and to run them
- report: this folder contains the source for this report

1.3 Usage

The entry point to the project is the **main.py** file. It can be called with the command line argument and provides two subcommands 'train' and 'run'. The first subcommand requires the name of the algorithm to be trained and stores a predictor in the path provided in the output argument, serialized using the Python pickle module.

```
1 $ python src/main.py train pegasos models/pegasos.pkl
```

The run subcommand takes a serialized model and then prints its training and test errors

```
1 $ python src/main.py run models/pegasos.pkl
```

There are other options available that are described using the '-help' option, they will be described in the next sections of this report as they come up.

2 Dataset analysis and preprocessing

2.1 Dataset description

The provided dataset contains 10000 points with 10 features named from x_1 to x_{10} and a label column named y .

All the feature are floating point values and the dataset is well formed (in the sense that there are no missing values).

The label column contains values that are either -1 or $+1$.

There isn't duplicated data in the training set.

I collect the major statistics from each feature in the dataset: mean, standard deviation, min and max

	x1	x2	x3	x4	x5
min	2.44342055e-03	-7.52493399e+00	9.85724553e+01	-7.07893888e+00	-9.99999717e-01
max	9.38422309e+00	8.30237476e+00	1.01260768e+02	-2.92150729e-06	9.99999998e-01
mean	1.59129826e+00	5.15879411e-01	9.98489361e+01	-1.50413876e+00	7.76447773e-02
std	1.32111881	2.05438485	0.71091203	1.13354878	0.70723419

	x6	x7	x8	x9	x10
min	-6.90697075e+00	-7.14075517e+00	-7.15188951e+00	-5.67739307e+01	-1.00000000e+00
max	8.76030588e+00	9.28726632e+00	6.21145227e+00	-5.42088897e+01	1.00000000e+00
mean	5.18228648e-02	9.75207134e-01	6.35194433e-01	5.19260973e-02	-5.54476783e+01
std	0.70471943	2.16212877	2.21259701	1.76955726	0.71004639

It's clear that the different features are not normalized and follow different probability distributions.

2.2 Feature Scaling

Because the dataset is already well formed and there are no missing values The first thing I do is scale the features.

This step should ensure that the values of the features are in a comparable range.

I tried two approaches: normalization and standardisation.

Scaling sets each feature to have a mean of 0 and a standard deviation of 1. This is achieved by subtracting the feature mean from each value and dividing by its standard deviation:

$$x' = \frac{x - \mu}{\sigma}$$

(where μ is the mean and σ is the standard deviation).

In the case of **normalization**, the features are rescaled in a fixed range between 0 and 1 in the following way:

$$x' = \frac{x - x_{min}}{x_{max} - x_{min}}$$

It's important to note that we should perform both approaches using a sound procedure, more specifically, we should avoid data leakage: When calculating the minimum and maximum of the feature, we should only consider the training set and scale the test set only according to these values, without deriving any information from it.

In the same way, we should calculate the mean and standard deviation for the standardisation process. Both of these approaches lead to important improvements and can be demonstrated both from an empirical point of view (see table TODO) and from a theoretical perspective.

2.3 Outliers removal

Another approach I tried is removing the outliers from the dataset using the Z-score methods. I calculated the score $Z = (x - \mu)/\sigma$ for each value where μ is the mean and σ is the variance of the feature.

I then removed all the points with a Z-score greater or equals than 3 in absolute value.

I found (and removed) 265 outliers (recall that the dataset has size 10000).

I tried training non-kernelized Perceptron and Pegasos over the modified dataset but the results shows that the dataset is already sufficiently cleaned: in fact it affects the performance of the models in a minimum way with no significative changes, and even in some cases it is (even if only slightly) worsening

	with outliers	without outliers
Perceptron	0.326	0.326
Pegasos	0.2865	0.29
Feature expanded Perceptron	0.087	0.087
Feature expanded Pegasos	0.0555	0.0565

Comparison of test error (using zero one loss) with and without outliers.

Note that when I refere to feature expansion I intend polynomial feature expansion of degree 2.

2.4 Feature correlation

I sorted all the point according to each axis and plot on the other axis all the other features to spot eventual correlation and I observed that the feature 2 and 5 have a linear correlation (with a negative coefficient, see Fig. 1) as the feature 5 and 9 (with positive coefficient, see Fig. 2).




Figure 1: Here there is a clear linear correlation between feature 2 and 5, but very noisy

One possibility in this case during the preprocessing of the data is to remove the correlated features and leave only one of them to avoid redundancy of the data.

I don't follow this approach because there is a sensible noise in the correlation and removing some




Figure 2: Another correlation is also present between feature 5 and 9 (with a negative coefficient)

features can lead to also removing this noise that can encode important information on the model.

2.5 Feature expansion

To being able to express non-homogeneous linear separators (hyperplane that don't pass through the origin) we add a constant feature of value 1 to each point in the dataset.

Let (\mathbf{x}) be any point in the dataset and (\mathbf{w}) be the linear separator, if we define $x' = (\mathbf{x}, 1)$ we can define $w' = (w, c)$, in that way:

$$w'^T x' = (\mathbf{w}^T \mathbf{x} + c)$$

3 Perceptron

The first algorithm that I implemented is the perceptron algorithm.

The perceptron algorithm is used to learn linear classifiers.

linear classifiers are identified by an hyperplane that separate the input space into two halfspaces, one positive and one negative.

The positive halfspace is called so because the dot product with the normal vector that identify the hyperplane and any point in that space is positive, similarly the negative halfspace has always negative dot products.

One important property of the Perceptron is the convergence (to the ERM) in a finite number of step if the dataset is linearly separable.

This properties is stated by the **Perceptron Convergence Theorem**:

Let $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$ be a linearly separable training set. Then the Perceptron algorithm returns a linear classifier with zero training error in a finite number of updates

$$M \leq \left(\min_{\mathbf{u}: \gamma(\mathbf{u}) \geq 1} \|\mathbf{u}\|^2 \right) \left(\max_{t=1, \dots, m} \|\mathbf{x}_t\|^2 \right)$$

where $\gamma(\mathbf{u})$ is the margin obtained by the linear separator \mathbf{u}

Is possible to show also a bound for non linearly separable cases:

$$M \leq \sum_{t=1}^T h_t(\mathbf{u}) + (\|\mathbf{u}\|X)^2 + \|\mathbf{u}\|X \sqrt{\sum_{t=1}^T h_t(\mathbf{u})} \quad \text{for all } \mathbf{u} \in \mathbb{R}^d$$

This shows a bound on the number of mistakes made by the Perceptron algorithm on any data sequence of arbitrary length T .

$h_t(\mathbf{u})$ is the hinge loss for the t -example.

Both the result show a linear dependence with the number of mistakes M and X^2 , the radius of the smaller sphere that inscribe all the training points.

This also show why, in our case, both *standardization* and *normalization* are so effective:

We are reducing the radius of this sphere and so having a tighter bound on the number of mistakes.

3.1 Naive version

Algorithm 1: The Perceptron algorithm

Input: Training set $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$

```
1  $\mathbf{w} = (0, \dots, 0)$ 
2 while true do
3   for  $i = 1, \dots, m$  do                                     // (epoch)
4     if  $y_i \mathbf{w}^\top \mathbf{x}_i \leq 0$  then
5        $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$                              // (update)
6     end
7   if no update in last epoch then
8     break
9 end
Output:  $\mathbf{w}$ 
```

My implementation slightly varies from the presented pseudocode: While the above on keep running until convergence if the training set is not lineary separable (as in this case) we never converge and so the algorithm never terminates.

To avoid this I use an additional parameter 'max_epoch' that limits the number of epoch which tha algorithm can run.

I choose to use a fixed value of 20 for this parameter both for the naive case and over the feature expanded dataset presented in the next section.

Training the perceptron algorithm with the preprocessing methodology described in the previous chapter (For all the algorihtms I describe I trained it with the default command line options), I have obtained a training error of 0.322625 and a test error of 0.326.

The linear separator founded by the perceptron algoihtm has the following features:

(0.55604295, 1.97486085, -2.58700382, -1.74490783, 1.91766823, -3.89876104, -0.02419966, 3.06371997, 0.21648717, -0.79054099, 1) (Recall the features are 11 because we add a constant feature of 1 to the dataset to being able to express non-homogeneous linear hyperplane).

3.2 Feature Expansion of 2nd degree

I also trained the perceptron algoihtm on a second degree polynomial feature expanded dataset.

In this version of the algorithm is possible to express hyperplane in a high dimensional feature space, and this can also be interpreted as a polynomial curve (of second degree in this case) in the original space.

For this reason the training and test error of the resulting predictor are significantly better than the previous version.

Specifically I obtained a training error of 0.085375, and a test error of 0.087.

This, of course, came at the cost of significantly increasing the number of features in the datasets and also the computational cost of operations between vectors (such as sums and dot products).

As we will see later, we can avoid this cost by using kernels, which allow us to compute the dot product of the feature expansion of two vectors without explicitly computing their expansion.

This are the features obtained:

(1.87559817e+01 1.13401244e+00 7.68047788e+00 -1.32276315e+01
1.73791816e+01 -5.35283741e+00 1.04738036e+01 6.65533633e+01
3.38354273e+01 5.20890557e+00 -1.40000000e+01 -2.37767269e+00
-1.15700558e+01 3.58887386e+00 -7.70998879e+00 -3.37225323e+00
8.62691037e+00 5.29455094e+00 7.16899813e+01 -2.75843723e+00
-1.44740075e+01 1.87559817e+01 1.58098829e+00 -2.16046998e+01

4.71738784e+00 -9.24568568e-01 2.31787163e+00 -6.88714236e-01
6.36558311e+00 2.24640419e+02 -3.01817078e+01 1.13401244e+00
-1.01581267e+01 -5.54650087e+00 -1.37119017e+01 5.92267047e+00
6.13176127e+00 -1.91940656e+01 1.20197686e+01 -3.78679497e+00
7.68047788e+00 -4.32369339e+00 -2.30745221e-01 -8.05450142e+00
1.43389046e+00 -7.36485907e+01 4.14731220e+00 9.17178493e-01
-1.32276315e+01 9.31783828e-01 -2.90129736e+01 1.10054038e+01
9.29221437e+00 5.16458889e-02 5.24165775e+00 1.73791816e+01
-3.39061351e+00 1.79056771e+01 1.68705706e+01 1.02044524e+01
-7.88104212e+00 -5.35283741e+00 4.04206126e+00 1.28786306e+01
1.01040343e+01 9.69582051e+00 1.04738036e+01 1.03988287e+01
2.78928092e+00 -2.37491314e+01 6.65533633e+01 8.51087821e-01
6.32907160e+00 3.38354273e+01 -2.72097857e+00 5.20890557e+00
-1.40000000e+01)

4 Support Vector Machine

The Support Vector Machine (SVM) algorithm learn linear classifiers, finding a linear classifier that is the **maximum margin separator hyperplane** and so achieve the maximum margin from all the point in the training set.

Given a linearly separable training set $S = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\} \in \mathbb{R}^d \times \{-1, 1\}$ it's possible to find this hyperplane solving the following convex optimization problem with linear constraints.

$$\begin{aligned} \min_{\mathbf{w} \in \mathbb{R}^d} \quad & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{s.t.} \quad & y_t \mathbf{w}^\top \mathbf{x}_t \geq 1 \text{ for } t = 1, \dots, m \end{aligned}$$

In case the training data is not separable, we try to minimise both how much of each constraint is violated and the margin of the separator.

We express this as another convex problem using the slack variables ξ_t and a regularisation coefficient λ . If the regularisation coefficient is large, the algorithms will generate a predictor that allows more classification error in the training set.

Conversely, if λ is small, we try to minimise the classification error we have made.

Usually for λ too small, we try to minimise the misclassification, so the training error is small and we are likely to overfit.

Instead, for choices of λ too large, we have a high training error, and if the test error is also high, we will underfit.

In the next subsection, I describe how I choose the hyperparameter of the regularisation coefficient and how the test and training errors vary when I change it more precisely.

$$\begin{aligned} \min_{(\mathbf{w}, \boldsymbol{\xi}) \in \mathbb{R}^{d+m}} \quad & \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{t=1}^m \xi_t \\ \text{s.t.} \quad & y_t \mathbf{w}^\top \mathbf{x}_t \geq 1 - \xi_t & t = 1, \dots, m \\ & \xi_t \geq 0 \text{ for } & t = 1, \dots, m \end{aligned}$$

Now, fix $\mathbf{w} \in \mathbb{R}^d$, we can see $\xi_t = [1 - y_t \mathbf{w}^\top \mathbf{x}_t]_+$ which is the hinge loss $h_t(\mathbf{w})$.

The SVM problem can be rewritten as

$$\min_{\mathbf{w} \in \mathbb{R}^d} \quad \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{t=1}^m h_t(\mathbf{w})$$

The optimization problem that describe the Support Vector Machine is optimized using the Pegasos algorithm.

The Pegasos algorithm is a variant of the Stochastic Gradient Descent algorithm, where at each step a point (or a set of points in the mini-batch variant) is sampled randomly from the training set and the current predictor is updated with the negative gradient of the loss of that training example weighted by a learning rate factor η_t .

In case of Pegasos the learning rate factor η_t is choose at each step as $\frac{1}{\lambda t}$.

I implement the Pegasos algorithm using the standard variant with the hinge loss, and with the logistic

loss (Described in the Logistic regression parameter).

Both this functions are convex upper bounds of the zero-one loss, and the λ regularization parameter allow to have a λ -strongly convex function to minimize with the gradient descent.

4.1 Naive

As I previously said I implemented Pegasos using two surrogate losses: hinge loss and logistic loss. Now I describe the implementation with hinge loss, that differs from the logistic one only by the update step.

Recall that the hinge loss is defined as $l(y, \hat{y}) = \max\{0, 1 - y_t \hat{y}\}$

Given $Z_t = (X_t, Y_t)$ a random sample from the training set, the update rule for Pegasos is:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla \ell_{Z_t}(\mathbf{w}_t)$$

Let be s_t the realization for the random variable Z_t

Where $\ell_{s_t}(\mathbf{w}) = [1 - y_{s_t} \mathbf{w}^T x_{s_t}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$ so

$$\nabla \ell_{s_t}(\mathbf{w}) = -y_{s_t} x_{s_t} \mathbb{I}\{h_{s_t}(\mathbf{w}) > 0\} + \lambda \mathbf{w}$$

Let $\mathbf{v}_t = y_t x_t \mathbb{I}\{h_t(\mathbf{w}_t) > 0\}$ and choosing $\eta_t = \frac{1}{\lambda t}$ we have

$$\mathbf{w}_{t+1} = \mathbf{w}_t \left(1 - \frac{1}{t}\right) + \frac{1}{\lambda t} \mathbf{v}_t$$

4.1.1 implementation details

I adapted this version of the algorithm from *Pegasos: Primal Estimated sub-GrAdient SOLver for SVM*[\[1\]](#)

Algorithm 2: Pegasos Algorithm

```

Input:  $S, \lambda, T$ 
1 for  $t = 1, 2, \dots, T$  do
2   Choose  $i_t$  uniformly at random.
3   Set  $\eta_t = \frac{1}{\lambda t}$ 
4   if  $y_{it} \mathbf{w}_t^T x_{i_t} < 1$  then
5     Set  $\mathbf{w}_{t+1} \leftarrow (1 - \eta_t \lambda) \mathbf{w}_t + \eta_t y_{it} \mathbf{x}_{i_t}$ 
6   end
7   else
8     Set  $\mathbf{w}_{t+1} \leftarrow (1 - \eta_t \lambda) \mathbf{w}_t$ 
9   end
10 end
11 Output  $\mathbf{w}_{T+1}$ 

```

The name of the variables are adapted to be consistent with the pseudo code reported

Between the code presented in the lecture and this one presented in the paper there are some differences with the pseudocode presented during the lectures:

- The gradient descent update is written in a slightly different way using a conditional statement instead of the classical indicator function, I choose to remain consistent also with this stylistic choice.

- Instead of return the average of all the weight vector calculated at each step, the paper returns only the last one.
The authors indicates that they notate an improvment in performance returning the last vector instead of the average.
I embrace also this variation.
- In the pseudo code of Pegasos they also describe an optional projection step to clamp the magnitude of the linear predictor, but I don't incorporate it.
- The author also provide a mini-batch version of the Pegasos algorithm, using the batch size k as another hyperparameter.

Another approach proposed by the paper is **sampling without replacement**: so a random permutation of the training set is choosen and the updates are performed in order on the new sequence of data. In this way, in one epoch, a training point is sampled only once.

After each epoch we can choose if we restart to sample data sequentially according to the same permutation or create a new one and sampling according that new order.

Although the authors report that this approaches gives better results than uniform sampling as I did, I haven't experiment this variant of the algorithm.

4.1.2 Hyperparameter tuning

Unlike the perceptron, the Pegasos algorithm had a hyperparameter to choose from: The regularisation coefficient λ .

To choose the best hyperparameter for this algorithm (and the others that I will implement), I choose the grid search method. I divide the training set into 2 different subsets: S_{train} and S_{dev} .

The validation set (S_{dev}) is used as a surrogate test set to get an estimate of the risk.

I choose a finite subset of the possible hyperparameter values ($\Theta_0 \subseteq \Theta$) and for each of them I create a predictor h_θ which is trained with the chosen hyperparameter on S_{train} .

The risk is then estimated using the validation error on each predictor, and the one with the lower risk is chosen.

The implementation of the grid search in my codebase is in the function 'grid_search', I used python know arguments to pass a dictionary where the key is the name of the hyperparameter (the same as used as an argument in the training algorithm) and the value is an iterator containing the subset of possible values (Θ_0).

In this case, I choose the set $\Theta_0 = \{0.0001, 0.001, 0.01, 0.1, 1, 10, 100\}$.

λ	S_{val}	S_{train}
0.0001	0.3725	0.3521666
0.001	0.2925	0.281333
0.01	0.273	0.27
0.1	0.267	0.2661666
1	0.273	0.269333
10	0.279	0.2775
100	0.2805	0.276333

The values S_{val} and S_{train} while λ changes.

Looking at the table we can deduce for which choice of λ we underfit and overfit and in our case the sweet spot is the value for $\lambda = 0.1$.

The test error for the best predictor is 0.2935.

We can also see that for smaller value of lambda we have an high test error and model is overfitting while for large value are likely to underfit.

The weight for the Pegasos algorithm are the following:

0.2080856 -0.0066276 0.07067274 -0.26477777 0.24795273 -0.06704699
0.28821758 0.65130349 0.16619769 -0.06650096 -0.0724

4.2 Logistic regression

Logistic regression aim to learn the function $\eta(\mathbf{x}) = \mathbb{P}(Y = +1 \mid \mathbf{X} = \mathbf{x})$.

The implementation of logistic regression differs from the standard Pegasos because it uses the logistic loss as surrogate loss and not hinge.

The logistic loss is defines ad follows:

$$\ell(y, \hat{y}) = \log_2(1 + e^{-y\hat{y}}).$$

So the gradient descent update becomes:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \eta_t \sigma(-y_t \mathbf{w}_t^T \mathbf{x}_t) y_t \mathbf{x}_t$$

I choose the regularization parameter in the same way and obtained the following result.

λ	S_{val}	S_{train}
0.1	0.278	0.2718333
1	0.2735	0.274333
10	0.28325	0.27816
100	0.2805	0.278166

The values S_{val} and S_{train} while λ changes for the logistic loss.

I have choosen the hyperparameter $\lambda = 1$ and obtained a test error of: 0.292.

4.3 Feature Expansion

I have also trained the two previous algorithm over the polynomial feature expansion of second degree of the whole dataset.

Like in the case of perceptron, since we are able to express a separator in an high dimension we obtain better results than the native version (at the cost of an increasing number of feature in the dataset).

4.3.1 Pegasos

This are the following validation and development errors training the Pegasos algorithm the same set of values choosen for the hyperparameter λ as in the naive version.

λ	S_{val}	S_{train}
0.001	0.045	0.039833
0.01	0.06	0.050166
0.1	0.0955	0.08833
1	0.1725	0.1525
10	0.219	0.211
100	0.256	0.246
1000	0.2635	0.25233

The test error obtained training the predictor on the whole training set with the regularization coefficient of 0.001 is 0.053 with a training error of 0.044125.

4.3.2 Logistic Regression

λ	S_{val}	S_{train}
0.1	0.11	0.1035
1	0.1675	0.15
10	0.197	0.188
100	0.2595	0.2485
1000	0.2505	0.2385

The test error obtained for the feature expanded logistic regression (with $\lambda = 0.001$) is 0.1125 with a training error of 0.1055.

Note that not all of the parameters for λ used in the previous case for the grid search are used in this case.

This is because for values that are too small, I get an overflow error when exponentiating in the sigmoid function.

These values, for which the gradient update calculation is not feasible, are simply discarded.

5 Kernel

In the previous section I have showed the benefit of training the algorithm in an high-dimensional space. This came with the cost of increasing sensibly the number of the feature even in case of a simple polynomial feature expansion of degree 2.

We can have the performance improvement obtained with feature expansion without increasing the dimensionality of the dataset using the kernel trick.

Instead of computing the feature expansion map on evry point and then computing the dot product with the expanded linear predictor we can use a kernel K that calculate the dot product of two expanded vector without explicitly calculate the expansion.

Let the feature expansion be $\phi(x)$ a kernel K for the feature expansion is defined as

$$K(x, x') = \phi(x)^T \phi(x')$$

for all

$$x, x' \in R$$

The two kind of kernel used in the project are the **Polynomial kernel** and the **Gaussian kernel**.

The polynomial kernel of degree n is

$$K_n(x, x') = (1 + x^T x')^n$$

The gaussian kernel of parameter γ is

$$K_\gamma(x, x') = \exp(-\frac{1}{2\gamma} \|x - x'\|^2)$$

References

- [1] Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, Andrew Cotter *Pegasos: Primal Estimated sub-GrAdient SOLver for SVM*. <https://home.ttic.edu/~nati/Publications/PegasosMPB.pdf>